

Extending JumpProcess.jl for fast point process simulation with time-varying intensities

Guilherme Augusto Zagatti¹, Samuel A. Isaacson³, Christopher Rackauckas⁴, Vasily Ilin⁵, See-Kiong Ng^{1,2}, and Stéphane Bressan^{1,2}

¹Institute of Data Science, National University of Singapore, Singapore

²School of Computing, National University of Singapore, Singapore

³Department of Mathematics and Statistics, Boston University

⁴Computer Science and AI Laboratory (CSAIL), Massachusetts Institute of Technology

⁵Department of Mathematics, University of Washington

ABSTRACT

Point processes model the occurrence of a countable number of random points over some support. They can model diverse phenomena, such as chemical reactions, stock market transactions and social interactions. We show that `JumpProcesses.jl` is a fast, general-purpose library for simulating point processes. `JumpProcesses.jl` was first developed for simulating jump processes via stochastic simulation algorithms (SSAs) (including Doob's method, Gillespie's methods, and Kinetic Monte Carlo methods). Historically, jump processes have been developed in the context of dynamical systems to describe dynamics with discrete jumps. In contrast, the development of point processes has been more focused on describing the occurrence of random events. In this paper, we bridge the gap between the treatment of point and jump process simulation. The algorithms previously included in `JumpProcesses.jl` can be mapped to three general methods developed in statistics for simulating evolutionary point processes. Our comparative exercise revealed that the library initially lacked an efficient algorithm for simulating processes with variable intensity rates. We, therefore, extended `JumpProcesses.jl` with a new simulation algorithm, `Coevolve`, that enables the rapid simulation of processes with locally-bounded variable intensity rates. It is now possible to efficiently simulate any point process on the real line with a non-negative, left-continuous, history-adapted and locally bounded intensity rate coupled or not with differential equations. This extension significantly improves the computational performance of `JumpProcesses.jl` when simulating such processes, enabling it to become one of the few readily available, fast, general-purpose libraries for simulating evolutionary point processes.

1. Introduction

Methods for simulating the trajectory of evolutionary point processes can be split into exact and inexact methods. Exact methods describe the realization of each point in the process chronologically. This exactness avoids bias from numerical approximations, but such methods can suffer from reduced performance when simulating systems with large populations (where numerous events can fire within a short period since every single point needs to be ac-

counted for). Inexact methods trade accuracy for speed by simulating the total number of events in successive intervals. They are popular in biochemical applications, e.g. τ -leap methods [4], which often require the simulation of chemical reactions in systems with large molecular populations.

Previously, point process simulation library development focused primarily on univariate processes with exotic intensities, or large systems with conditionally constant intensities, but not on both. As such, there was no widely used general-purpose software for efficiently simulating compound point processes in large systems with time-dependent rates. To enable the efficient simulation of such processes, we contribute the `Coevolve` aggregator to `JumpProcesses.jl`, a core component of the popular `DifferentialEquations.jl` library [17]. The implemented algorithm improves the `COEVOLVE` algorithm described in [2] from where it borrows its name. Among other improvements, our algorithm supports any process with locally bounded conditional intensity rates, adapts to intensity rates that can change between jumps, can be coupled with differential equations, and avoids both the unnecessary re-computation of randomly generated numbers and the computation of the intensity rate when its lower bound is available. This extension of `JumpProcesses.jl` dramatically boosts the computational performance of the library in simulating processes with intensities that have an explicit dependence on time and/or other continuous variables, significantly expanding the type of models that can be efficiently simulated. Widely-used point processes with such intensities include compound inhomogeneous Poisson, Hawkes, and stress-release processes — all described in [1]. Since `JumpProcesses.jl` is a member of Julia's SciML organization, it also becomes easier, and more feasible, to incorporate compound point processes with explicit time-dependent rates into a wide variety of applications and higher-level analyses. With our new additions we bump `JumpProcesses.jl` to version 9.7¹.

In this paper, we bridge the gap between simulation methods developed in statistics and biochemistry, which led us to the development of `Coevolve`. First, we briefly introduce evolutionary point processes. Next, since all simulation methods require a basic under-

¹All examples and benchmarks in this paper use this version of the library

77 standing of simulation methods for the Poisson homogeneous process, we first describe such methods. Then, we identify and discuss
 78 three general, exact methods. In the second part of this paper, we
 79 describe the algorithms in `JumpProcesses.jl` and how they relate
 80 to the literature. We highlight our contribution `Coevolve`, investigate
 81 the correctness of our implementation and provide performance
 82 benchmarks to demonstrate its value. The paper concludes
 83 by discussing potential improvements.
 84

85 2. The evolutionary point process

86 The evolutionary point process is a stochastic collection of marked
 87 points over a one-dimensional support. They are exhaustively described
 88 in [1]. The likelihood of any evolutionary point process is fully
 89 characterized by its conditional intensity,

$$\lambda^*(t) \equiv \lambda(t | H_{t^-}) = \frac{p^*(t)}{1 - \int_{t^-}^{t_n} p^*(u) du}, \quad (2.1)$$

90 and conditional mark distribution, $f^*(k|t)$ — see Chapter 7 [1].
 91 Here $H_{t^-} = \{(t_n, k_n) \mid 0 \leq t_n \leq t\}$ denotes the internal
 92 history of the process up to but not including t , the superscript $*$
 93 denotes the conditioning of any function on H_{t^-} , and $p^*(t)$ is the
 94 density function corresponding to the probability of an event taking
 95 place at time t given H_{t^-} . We can interpret the conditional intensity
 96 as the likelihood of observing a point in the next infinitesimal
 97 unit of time, given that no point has occurred since the last observed
 98 point in H_{t^-} . Lastly, the mark distribution denotes the density function
 99 corresponding to the probability of observing mark k given the
 100 occurrence of an event at time t and internal history H_{t^-} .

101 3. The homogeneous process

102 A homogeneous process can be simulated using properties of the
 103 Poisson process, which allow us to describe two equivalent sampling
 104 procedures. The first procedure consists of drawing successive
 105 inter-arrival times. The distance between any two points in a
 106 homogeneous process is distributed according to the exponential
 107 distribution — see Theorem 7.2 [9]. Given the homogeneous
 108 process with intensity λ , then the distance Δt between two points
 109 is distributed according to $\Delta t \sim \exp(\lambda)$. Draws from the
 110 exponential distribution can be performed by drawing from a uniform
 111 distribution in the interval $[0, 1]$. If $V \sim U[0, 1]$, then
 112 $T = -\ln(V)/\lambda \sim \exp(1)$. (Note, however, in Julia the optimized
 113 Zigurat-based method used in the `randexp` stdlib function is generally
 114 faster than this *inverse* method for sampling a unit exponential
 115 random variable.) When a point process is homogeneous, the *inverse*
 116 method of Subsection 4.1 reduces to this approach. Thus, we defer
 117 the presentation of this Algorithm to the next section.
 118

119 The second procedure uses the fact that Poisson processes can be
 120 represented as a mixed binomial process with a Poisson mixing
 121 distribution — see Proposition 3.5 [9]. In particular, the total number
 122 of points of a Poisson homogeneous process in $[0, T)$ is distributed
 123 according to $\mathcal{N}(T) \sim \text{Poisson}(\lambda T)$ and the location of each point
 124 within the region is distributed according to the uniform distribution
 125 $t_n \sim U[0, T]$.

126 4. Exact simulation methods

127 4.1 Inverse methods

128 The *inverse* method leverages Theorem 7.4.I [1] which states that
 129 every simple point process² can be transformed to a homogeneous
 130 Poisson process with unit rate via the compensator. Let t_n be the
 131 time in which the n -th chronologically sorted event took place and
 132 $t_0 \equiv 0$, we define the compensator as:

$$\Lambda^*(t_n) \equiv \tilde{t}_n \equiv \int_0^{t_n} \lambda^*(u) du \quad (4.1)$$

133 The transformed data \tilde{t}_n forms a homogeneous Poisson process
 134 with unit rate. Now, if this is the case, then the transformed interval
 135 is distributed according to the exponential distribution.

$$\Delta \tilde{t}_n \equiv \tilde{t}_n - \tilde{t}_{n-1} \sim \exp(1) \quad (4.2)$$

136 The idea is to draw realizations from the unit rate Exponential process
 137 and solve Equation 4.2 for t_n to determine the next event/firing
 138 time. We illustrate this in Algorithm 1 where we adapt Algorithm
 139 7.4 [1].

140 Whenever the conditional intensity is constant between two
 141 points, Equation 4.2 can be solved analytically. Let $\lambda^*(t) =$
 142 $\lambda_{n-1}, \forall t_{n-1} \leq t < t_n$, then

$$\begin{aligned} \int_{t_{n-1}}^{t_n} \lambda^*(u) du &= \Delta \tilde{t}_n \iff \\ \lambda_{n-1}(t_n - t_{n-1}) &= \Delta \tilde{t}_n \iff \\ t_n &= t_{n-1} + \frac{\Delta \tilde{t}_n}{\lambda_{n-1}}. \end{aligned} \quad (4.3)$$

143 Which is equivalent to drawing the next realization time from the
 144 re-scaled exponential distribution $\Delta t_n \sim \exp(\lambda_{n-1})$. As we will
 145 see in Subsection 2, this implies that the *inverse* and *thinning*
 146 methods are the same whenever the conditional intensity is constant
 147 between jumps.

148 The main drawback of the *inverse* method is that the root finding
 149 problem defined in Equation 4.2 often requires a numerical solution.
 150 To get around a similar obstacle in the context of the piecewise
 151 deterministic Markov process, Veltz [23] proposes a change of
 152 variables in time that recasts the root finding problem into an
 153 initial value problem. He denotes his method *CHV*.

154 Piecewise deterministic Markov processes are composed of two
 155 parts: the jump process and the piecewise ODE that changes
 156 stochastically at jump times — see Lemaire *et al.* [11] for a formal
 157 definition. Therefore, it is easy to employ *CHV* in our case
 158 by setting the ODE part to zero throughout time. Adapting from
 159 Veltz [23], we can determine the model jump time t_n after sampling
 160 $\Delta \tilde{t}_n \sim \exp(1)$ by solving the following initial value problem
 161 until $\Delta \tilde{t}_n$.

$$t(0) = t_{n-1}, \quad \frac{dt}{dt} = \frac{1}{\lambda^*(t)} \quad (4.4)$$

162 Looking back at Equation 4.1, we note that it is a one-to-one mapping
 163 between t and \tilde{t} which makes it completely natural to write
 164 $t(\Delta \tilde{t}_n) \equiv \Lambda^{*-1}(\tilde{t}_{n-1} + \Delta \tilde{t}_n)$.

²A simple point process is a process in which the probability of observing more than one point in the same location is zero.

165 Alternatively, when the intensity function is differentiable between 203
 166 jumps we can go even further by recasting the jump problem as a 204
 167 piecewise deterministic Markov process. Let $\lambda_n^* \equiv \lambda^*(t_n)$, then 205
 168 the flow $\varphi_{t-t_n}(\lambda_n^*)$ maps the initial value of the conditional intensity 206
 169 at time t_n to its value at time t . In other words, the flow describes 207
 170 the deterministic evolution of the conditional intensity function 208
 171 over time. Next, denote $\mathbf{1}(\cdot)$ as the indicator function, then the 209
 172 conditional intensity function can be re-written as a jump process: 210

$$\lambda^*(t) = \sum_{n \geq 1} \varphi_{t-t_{n-1}}(\lambda_{n-1}^*) \mathbf{1}(t_{n-1} \leq t < t_n). \quad (4.5)$$

173 According to Meiss [15], if $\varphi_t(\cdot)$ is a flow, then it is a solution to 211
 174 the initial value problem: 212

$$\varphi_0(\lambda_n^*) = \lambda_n^*, \quad \frac{d}{dt} \varphi_{t-t_n}(\lambda_n^*) = g(\varphi_{t-t_n}(\lambda_n^*)) \quad (4.6)$$

175 where $g: \mathbb{R}^+ \rightarrow \mathbb{R}$ is the vector field of λ^* such that $d\lambda^*/dt =$ 213
 176 $g(\lambda^*)$. 214

177 Based on Equation 2.1, we find that the probability of observing an 215
 178 interval longer than s given internal history H_{t-} is equivalent to:

$$\begin{aligned} \Pr(t_n - t_{n-1} > s \mid H_{t-}) &= 1 - \int_{t_{n-1}}^{t_{n-1}+s} p^*(u) du = \\ &= \exp\left(-\int_{t_{n-1}}^{t_{n-1}+s} \lambda^*(u) du\right) = \\ &= \exp\left(-\int_{t_{n-1}}^{t_{n-1}+s} \varphi_{u-t_{n-1}}(\lambda_{n-1}^*) du\right) \end{aligned} \quad (4.7)$$

179 Equations 4.5 and 4.7 define a piecewise deterministic Markov process 216
 180 satisfying the conditions of Theorem 3.1 [23]. In this case, we 217
 181 find t_n by solving the following initial value problem from 0 to 218
 182 $\Delta \tilde{t}_n \sim \exp(1)$. 219

$$\begin{cases} \lambda^*(t(0)) = \lambda^*(t_{n-1}), \quad \frac{d\lambda^*}{dt} = \frac{g(\lambda^*(t))}{\lambda^*(t)} \\ t(0) = t_{n-1}, \quad \frac{dt}{dt} = \frac{1}{\lambda^*(t)}. \end{cases} \quad (4.8)$$

183 This problem specifies how the conditional intensity and model 220
 184 time evolve with respect to the transformed time. The solution to 221
 185 Equation 4.2 is then given by $(t_n = t(\Delta \tilde{t}_n), \lambda^*(t(\Delta \tilde{t}_n))) =$ 222
 186 $\lambda^*(t_n)$. 223

187 In Algorithm 1, we can implement the CHV method by solving 224
 188 either Equation 4.4 or Equation 4.8 instead of Equation 4.2. We 225
 189 denote the first specification as *CHV simple* and the second as 226
 190 *CHV full*. Note that *CHV full* requires that the conditional intensity 227
 191 be piecewise differentiable. The algorithmic complexity is then 228
 192 determined by the ODE solver and no root-finding is required. In 229
 193 Section 6.2, we will show that there are substantial differences in 230
 194 performance between them with the full specification being faster. 231
 195 Another concern with Algorithm 1 is updating and drawing from 232
 196 the conditional mark distribution in Line 8, and updating the conditional 233
 197 intensity in Line 9. Assume a process with K number of 234
 198 marks. A naive implementation of Line 9 scales with the number 235
 199 of marks as $O(K)$ since λ^* is usually constructed as the sum of K 236
 200 independent processes, each of which requires updating the conditional 237
 201 intensity rate. Likewise, drawing from the mark distribution 238
 202 in Line 8 usually involves drawing from a categorical distribution 239

203 whose naive implementations also scales with the number of marks 204
 205 as $O(K)$. 206

207 Finally, Algorithm 1 is not guaranteed to terminate in finite time 208
 209 since one might need to sample many points before $t_n > T$. The 210
 211 sampling rate can be especially high when simulating the process 212
 213 in a large population with self-exciting encounters. In biochemistry, 214
 215 Salis and Kaznessis [19] partition a large system of chemical reactions 216
 217 into two: fast and slow reactions. While they approximate the fast 218
 219 reactions with a Gaussian process, the slow reactions are solved 219
 220 using a variation of the inverse method. They obtain an equivalent 220
 221 expression for the rate of slow reactions as in Equation 4.2, which 221
 222 is integrated with the Euler method. 222

Algorithm 1 The *inverse* method for simulating a marked evolutionary point process over a fixed duration of time $[0, T)$.

```

1: procedure INVERSEMETHOD( $[0, T)$ ,  $\lambda^*$ ,  $f^*$ )
2:   initialize the history  $H_{T-} \leftarrow \{\}$ 
3:   set  $n \leftarrow 0, t \leftarrow 0$ 
4:   while  $t < T$  do
5:      $n \leftarrow n + 1$ 
6:     draw  $\Delta \tilde{t}_n \sim \exp(1)$ 
7:     find the next event time  $t_n$  by solving Equation 4.2 or 4.8
8:     update  $f^*$  and draw the mark  $k_n \sim f^*(k \mid t_n)$ 
9:     update the history  $H_{T-} \leftarrow H_{T-} \cup (t_n, k_n)$  and  $\lambda^*$ 
10:  end while
11:  return  $H_{T-}$ 
12: end procedure

```

4.2 Thinning methods

Thinning methods are one of the most popular methods for simulating point processes. The main idea is to successively sample a homogeneous process, then thin the obtained points with the conditional intensity of the original process. As stated in Proposition 7.5.I [1], this procedure simulates the target process by construction. The advantage of *thinning* over *inverse* methods is that the former only requires the evaluation of the conditional intensity function while the latter requires computing the inverse of its integrated form [1].

Thinning algorithms have been proposed in different forms [1]. The Shedler-Lewis algorithm can simulate processes with bounded intensity [12]. The classical algorithm from Ogata [16] overcomes this limitation and only requires the local boundedness of the conditional intensity. The advantage of Ogata's algorithm and its variations is that it can simulate processes with potentially unbounded intensity, such as self-exciting ones. As long as the intensity conditioned on the simulated history remains locally bounded, it is possible to simulate subsequent points indefinitely.

In biochemistry, the *thinning* method was popularized by Gillespie [6, 5]. For this reason, this method is also called the *Gillespie* method. Gillespie himself called it the *direct* method or the *stochastic simulation algorithm*. Gillespie introduced the *thinning* method in the context of simulating chemical reactions of well-stirred systems. He developed a stochastic model for molecule interactions from physics principles without any references to the point process theory developed in this section. His model of chemical interactions is equivalent to a marked Poisson process with constant conditional intensity between jumps. The model consists of distinct populations of molecular species that interact through several reaction channels. A chemical reaction consists of a Poisson process that transforms a set of molecules of some type into a set

247 of molecules of another type. What Gillespie calls the master equation can be deduced from the *superposition theorem* — Theorem 3.3 [9].

250 Alternatively, in biochemistry, *thinning* methods are known as *rejection* algorithms. Than *et al.* [21, 22] proposed the *rejection-based algorithm with composition-rejection search*, yet another more sophisticated variation of the *thinning* method. In this case, the procedure groups similar processes together. For each group, an upper- and lower-bound conditional intensity is used for thinning. A similar procedure is also described in [20], in which the authors refer to their algorithm as *kinetic Monte Carlo*.

258 In Algorithm 2, we modify Algorithm 7.5.IV [1] to incorporate the idea of a lower bound on the conditional intensity from [22]. To implement the algorithm, we define three functions, $\bar{B}^*(t) = \bar{B}(t | H_t)$, $\underline{B}^*(t) = \underline{B}(t | H_t)$ and $L^*(t) = L(t | H_t)$, that characterize the local boundedness condition such that:

$$\lambda^*(t+u) \leq \bar{B}^*(t) \text{ and } \lambda^*(t+u) \geq \underline{B}^*(t), \quad (4.9)$$

$$\forall 0 \leq u \leq L^*(t).$$

263 The tighter the bound on $\bar{B}^*(t)$, the lower the number of samples discarded. Since looser bounds lead to less efficient algorithms, the art, when simulating via *thinning*, is to find the optimal balance between the local supremum of the conditional intensity $\bar{B}^*(t)$ and the duration of the local interval $L^*(t)$. On the other hand, the infimum $\underline{B}^*(t)$ can be used to avoid the evaluation of $\lambda^*(t+u)$ in Line 5 of Algorithm 3 which often can be expensive.

270 When the conditional intensity is constant between jumps such that $\lambda^*(t) = \lambda_{n-1}, \forall t_{n-1} \leq t < t_n$, let $\bar{B}^*(t) = \underline{B}^*(t) = \lambda_{n-1}$ and $L^*(t) = \infty$. We have that for any $u \sim \exp(1 / \bar{B}^*(t)) = \exp(\lambda_{n-1})$ and $v \sim U[0, 1], u < L^*(t) = \infty$ and $v < \lambda^*(t+u) / \bar{B}^*(t) = 1$. Therefore, we advance the internal history for every iteration of Algorithm 2. In this case, the bound $\bar{B}^*(t)$ is as tight as possible, and this method becomes the same as the *inverse* method of Subsection 4.1.

278 We can draw many more connections between the *thinning* and *inverse* methods. Lemaire *et al.* [11] propose a version of the *thinning* algorithm for Piecewise Deterministic Markov Processes which does not use the local interval L^* for rejection — this is equivalent to $L^*(t) = \infty$ —, and does not assume the upper bound $\bar{B}^*(t)$ is constant over $L^*(t)$. The efficiency of their algorithm depends on the assumption that the stochastic process determined by $\bar{B}^*(t)$ can be efficiently inverted such that candidate times can be efficiently obtained using Equation 4.1. They propose an optimal bound as a piecewise constant function partitioned in such a way that it envelopes the intensity function as strictly as possible. They then show that under certain conditions the stochastic process determined by $\bar{B}^*(t)$ converges in distribution to the target conditional intensity as the partitions of the optimal boundary converge to zero. Although their simulation approach does not exactly match ours, it suggests some properties between the *thinning* and the *inverse* method that we could investigate in the future. Among other things, the efficiency of *thinning* compared to *inversion* most likely depends on the rejection rate obtained by the former and the number of steps required by the ODE solver for the latter.

298 While *thinning* algorithms avoid the issue of directly computing the inverse of the integrated conditional intensity, they increase the number of time steps needed in the sampling algorithm as we are now sampling from a process with an increased intensity relative to the original process. Moreover, like the *inverse* method, *thinning* algorithms can also face issues related with drawing from the conditional mark distribution — Line 11 of Algorithm 2 —, and

updating the conditional intensity — Line 3 of Algorithm 3 — and the mark distribution — Line 12 of Algorithm 2.

Algorithm 2 The *thinning* method for simulating a marked evolutionary point process over a fixed duration of time $[0, T)$.

```

1: procedure THINNINGMETHOD( $[0, T), \lambda^*, f^*$ )
2:   initialize the history  $H_{T^-} \leftarrow \{\}$ 
3:   set  $n \leftarrow 0, t \leftarrow 0$ 
4:   while true do
5:      $t \leftarrow \text{TimeViaThinning}([t, T), H_{T^-}, \lambda^*)$ 
6:     if  $t \geq T$  then
7:       break
8:     end if
9:      $n \leftarrow n + 1$ 
10:     $t_n \leftarrow t$ 
11:    update  $f^*$  and draw the mark  $k_n \sim f^*(k | t_n)$ 
12:    update the history  $H_{T^-} \leftarrow H_{T^-} \cup (t_n, k_n)$ 
13:  end while
14:  return  $H_{T^-}$ 
15: end procedure

```

Algorithm 3 Generates the next event time via *thinning*.

```

1: procedure TIMEVIATHINNING( $[t, T), \lambda^*, H_t$ )
2:   while  $t < T$  do
3:     update  $\lambda^*$ 
4:     find  $\bar{B}^*(t), \underline{B}^*(t)$  and  $L^*(t)$  which satisfy Eq. 4.9
5:     draw  $u \sim \exp(\bar{B}^*(t))$  and  $v \sim U[0, \bar{B}^*(t)]$ 
6:     if  $u > L^*(t)$  then
7:        $t \leftarrow t + L^*(t)$ 
8:     next
9:     end if
10:    if  $(v > \underline{B}^*(t))$  and  $(v > \lambda^*(t+u))$  then
11:       $t \leftarrow t + u$ 
12:    next
13:    end if
14:     $t \leftarrow t + u$ 
15:    break
16:  end while
17:  return  $t$ 
18: end procedure

```

4.3 Queuing methods

307 As an alternative to his *direct* method — in this text referred as the constant rate *thinning* method —, Gillespie introduced the *first reaction* method in his seminal work on simulation algorithms [6]. The *first reaction* method separately simulates the next reaction time for each reaction channel — *i.e.* for each mark. It then selects the smallest time as the time of the next event, followed by updating the conditional intensity of all channels accordingly. This is a variation of the constant rate *thinning* method to simulate a set of inter-dependent point processes, making use of the *superposition theorem* — Theorem 3.3 [9] — in the inverse direction. Gibson and Bruck [3] improved the *first reaction* method with the *next reaction* method. They innovate on three fronts. First, they keep a priority queue to quickly retrieve the next event. Second, they keep a dependency graph to quickly locate all conditional intensity rates that need to be updated after an event is fired. Third,

they re-use previously sampled reaction times to update unused reaction times. This minimizes random number generation, which can be costly. Priority queues and dependency graphs have also been used in the context of social media [2] and epidemics [8] simulation. In both cases, the phenomena are modelled as point processes.

We prefer to call this class of methods *queuing* methods since most efficiency gains come from maintaining a priority queue of the next event times. Up to this point we assumed that we were sampling from a global process with a mark distribution that could generate any mark k given an event at time t . With queuing, it is possible to simulate point processes with a finite space of marks as M interdependent point processes — see Definition 6.4.1 [1] of multivariate point processes — doing away with the need to draw from the mark distribution at every event occurrence. Alternatively, it is possible to split the global process into M interdependent processes each one of which with its own mark distribution.

Our contribution, Algorithm 5, presents a method for sampling a superposed point process consisting of M processes by keeping the strike time of each process in a priority queue Q . The priority queue is initially constructed in $O(M)$ steps in Lines 4 to 7 of Algorithm 5. In contrast to *thinning* methods, updates to the conditional intensity depend only on the size of the neighborhood of i . That is, processes j whose conditional intensity depends on the history of i . If the graph is sparse, then updates will be faster than with *thinning*.

A source of inefficiency in some implementations of queuing algorithms is the fact that one might need to go through multiple rejection cycles before accepting a time candidate t_i for process i . This might require looking ahead in the future. In addition to that, if process j , which i depends on, takes place before i , then we need to repeat the whole thinning process to obtain a new time candidate for i . We thus propose Algorithm 5 which is a queuing algorithm that performs thinning in synchrony with the main loop, thus avoiding look ahead and wasted rejections. Since thinning is now synced with the main loop, it is possible to couple this simulator with other algorithms that step chronologically through time. These include ordinary differential equation solvers, enabling us to simulate jump processes with rates given by a differential equation. This is the first synced thinning algorithm we are aware of.

Algorithm 4 Generates the next candidate time for *queuing*.

```

1: procedure QUEUE TIME( $t, \lambda^*, H_t$ )
2:   update  $\lambda^*$ 
3:   find  $\bar{B}^*(t), \underline{B}^*(t)$  and  $L^*(t)$  which satisfy Eq. 4.9
4:   draw  $u \sim \exp(\bar{B}^*(t))$ 
5:   if  $u > L^*(t)$  then
6:     accepted  $\leftarrow$  false
7:   else
8:     accepted  $\leftarrow$  true
9:   end if
10:   $t \leftarrow t + u$ 
11:  return  $t, \bar{B}^*(t), \underline{B}^*$ , accepted
12: end procedure

```

5. Implementation

JumpProcesses.jl is a Julia library for simulating jump — or point — processes which is part of Julia’s SciML organization. Our discussion in Section 4 identified three exact methods for simulating point processes. In all the cases, we identi-

Algorithm 5 The *queuing* method for simulating a marked evolutionary point process over a fixed duration of time $[0, T)$.

```

1: procedure QUEUING METHOD( $[0, T), \{\lambda_k^*\}, \{f_k^*\}$ )
2:   initialize the history  $H_{T^-} \leftarrow \{\}$ 
3:   set  $n \leftarrow 0, t \leftarrow 0$ 
4:   for  $i=1, M$  do
5:      $(t_i, \bar{B}_i^*, \underline{B}_i^*, a_i) \leftarrow$  QueueTime( $0, H_{T^-}, \lambda_i^*(\cdot)$ )
6:     push  $(t_i, \bar{B}_i^*, \underline{B}_i^*, a_i, i)$  to  $Q$ 
7:   end for
8:   while  $t < T$  do
9:     first  $(t_i, i, \bar{B}_i^*, \underline{B}_i^*, a_i, i)$  from  $Q$ 
10:     $t \leftarrow t_i$ 
11:    if  $t \geq T$  then
12:      break
13:    end if
14:    draw  $v \sim U[0, \bar{B}_i^*]$ 
15:    if  $(v > \underline{B}_i^*)$  and  $(v > \lambda^*(t))$  then
16:       $a_i \leftarrow$  false
17:    end if
18:    if  $a_i$  then
19:       $n \leftarrow n + 1$ 
20:       $t_n \leftarrow t$ 
21:      update  $f^*$  and draw the mark  $k_n \sim f_i^*(k | t_n)$ 
22:      update the history  $H_{T^-} \leftarrow H_{T^-} \cup (t_n, k_n)$ 
23:      for  $j \in \{i\} \cup \text{Neighborhood}(i)$  do
24:         $(t_j, \bar{B}_j^*, \underline{B}_j^*, a_j) \leftarrow$  QueueTime( $0, H_{T^-}, \lambda_j^*(\cdot)$ )
25:        update  $(t_j, \bar{B}_j^*, \underline{B}_j^*, a_j, j)$  in  $Q$ 
26:      end for
27:    else
28:       $(t_i, \bar{B}_i^*, \underline{B}_i^*, a_i) \leftarrow$  QueueTime( $0, H_{T^-}, \lambda_i^*(\cdot)$ )
29:      update  $(t_i, \bar{B}_i^*, \underline{B}_i^*, a_i, i)$  in  $Q$ 
30:    end if
31:  end while
32:  return  $H_{T^-}$ 
33: end procedure

```

fied two mathematical constructs required for simulation: the intensity rate and the mark distribution. In JumpProcesses.jl, these can be mapped to user defined functions rate(u, p, t) and affect!(integrator). The library provides APIs for defining processes based on the nature of the intensity rate and the intended simulation algorithm. Processes intended for exact methods can choose between ConstantRateJump and VariableRateJump. While the former expects the rate between jumps to be constant, the latter allows for time-dependent rates. The library also provides the MassActionJump API to define large systems of point processes that can be expressed as reaction equations. Finally, RegularJump are intended for inexact methods.

The *inverse* method as described around Equation 4.2 uses root find to find the next jump time. Jumps to be simulated via the *inverse* method must be initialized as a VariableRateJump. JumpProcesses.jl builds a continuous callback following the algorithm in [19] and passes the problem to an OrdinaryDiffEq.jl integrator, which easily interoperates with JumpProcesses.jl (both libraries are part of the SciML organization, and by design built to easily compose). JumpProcesses.jl does not yet support the CHV ODE based approach.

Alternatively, *thinning* and *queuing* methods can be simulated via discrete steps. In the context of the library, any method that uses a discrete callback is called an *aggregator*. There are twelve differ-

ent aggregators, seven of which implement a variation of the *thinning* method and five of which a variation of the *queuing* method. We start with the *thinning* aggregators, none of which support `VariableRateJump`. Algorithm 2 assumes that there is a single process. In reality, all the implementations assume a finite multivariate point process with M interdependent processes. This can be easily conciliated, as we do now, using Definition 6.4.1 [1] which states the equivalence of such process with a point process with a finite space of marks. As all the *thinning* aggregators only deal with `ConstantRateJump`, the intensity between jumps is constant. Algorithm 3 short-circuits to quickly return $t \sim \exp(\bar{B}) = \exp(\lambda_n)$ as discussed in Subsection 4.2. Next, the mark distribution becomes the categorical distribution weighted by the intensity of each process. That is, given an event at time t_n , we have that the probability of drawing process i out of M sub-processes is $\lambda_i^*(t_n)/\lambda^*(t_n)$. Conditional on sub-process i , the corresponding `affect!(integrator)` is invoked, that is, $k_n \sim f_i^*(k | t_n)$. Here we use a notation analogous to Section 4.3.

Where most implementations differ is on updating the mark distribution in Line 11 of Algorithm 2 and the conditional intensity rate in Line 3 of Algorithm 3. `Direct` and `DirectFW` follows the *direct* method in [6] which re-evaluates all intensities after every iteration scaling at $O(K)$. When drawing the process to fire, it executes a search in an array that stores the cumulative sum of rates. `DirectCR`, `SortingDirect` and `RDirect` only re-evaluate the intensities of the processes that are affected by the realized process. This operation is executed efficiently by keeping a vector of dependencies. These three algorithms differ in how they select the process. `DirectCR` keeps the intensity rates in a priority table, it is implemented after [20]. `SortingDirect` keeps the intensity rate in a loosely sorted array following [14]. In both cases, the idea is to use a randomly generated number between zero and one to guide the search for the next jump. With the intensity rates sorted, more frequent processes should be selected faster than less frequent ones. Overall, this should increase the speed of the simulation. `RDirect` keeps track of the maximum rate of the system, it implements an algorithm equivalent to *thinning* with \bar{B} equal to the maximum rate. However, the implementation differs. It thins with $\bar{B} = \lambda_n$, then randomly selects a candidate process and confirms the candidate only if its rate is above a random proportion of the maximum rate. Finally, `RSSA` and `RSSACR` group processes with similar rates in bounded brackets. The upper bounds are used for *thinning*. For each round of *thinning*, a sampled candidate process is considered for selection. In `RSSA`, the candidate process is selected similarly to `Direct`, while a priority table is used in `RSSACR`. Both of these algorithms follow from [21, 22].

Next, we consider the *queuing* aggregators. Starting with aggregators that only support `ConstantRateJumps` we have, `FRM`, `FRMFW` and `NRM`. `FRM` and `FRMFW` follow the *first reaction* method in [6]. To compute the next jump, both algorithms compute the time to the next event for each process and select the process with minimum time. This is equivalent to assuming a complete dependency graph in Algorithm 5. For large systems, they can be less efficient than `NRM`. The latter implementation is sourced from [3] and follows Algorithm 5 very closely.

Previously, we attempted to bridge the gap between the treatment of point process simulation in statistics and biochemistry. Despite the many commonalities, most of the algorithms implemented in `JumpProcesses.jl` are derived from the biochemistry literature. There has been less emphasis on implementing processes commonly studied in statistics such as self-exciting point processes characterized by time-varying and history-dependent inten-

sity rates. This is addressed by our latest aggregator, `Coevolve`. This is the first aggregator that supports `VariableRateJumps`, facilitating substantially more performant simulation of processes with time-dependent intensity rates in `JumpProcesses.jl` and `DifferentialEquations.jl` compared to the current *inverse* method-based approach that relies on ODE integration and continuous events.

The implementation of this aggregator takes inspiration from [2], and improves the method in several ways. First, we take advantage of the modularity and composability of Julia to design an API that accepts any intensity rate, not only the Hawkes'. Second, we avoid the re-computation of unused random numbers. When updating processes that have not yet fired, we can transform the unused time of constant rate processes to obtain the next candidate time for the first round of iteration of the *thinning* procedure in Algorithm 3. This saves one round of sampling from the exponential distribution, which translates into a faster algorithm. Third, we allow the user to supply a lower bound rate which can short-circuit the loop in Algorithm 3, saving yet another round of sampling. Fourth, it adapts to processes with constant intensity between jumps which reduces the loop in Algorithm 3 to the equivalent implemented in `NRM`. Finally, since `Coevolve` can be mapped to a *thinning* algorithm — see [2] —, it can simulate any point process on the real line with a non-negative, left-continuous, history-adapted and locally bounded intensity rate as per Proposition 7.5.1 [1].

`Coevolve` syncs with the main execution loop which means that it can be easily coupled with differential equations modeled with `OrdinaryDiffEq.jl`. For instance, It is possible to model processes whose rates are given by a differential equation. This is a departure from the algorithm described in [2] which translates not only into a faster, but also more flexible simulator. This difference in implementation follows our discussion on the relationship between the main execution loop and the thinning loop in Section 4.3.

6. Empirical evaluation

This section conducts some empirical evaluation of the `JumpProcesses.jl` aggregators described in Section 5. First, since `Coevolve` is a new aggregator, we test its correctness by conducting statistical analysis. Second, we conduct the jump benchmarks available in `SciMLBenchmarks.jl`. We have added new benchmarks that assess the performance of the new aggregators under settings that could not be simulated with previous aggregators.

6.1 Statistical analysis of `Coevolve`

To simulate a process intended for a discrete solver with `JumpProcesses.jl`, we define a discrete problem, initialize the jumps and define the jump problem which takes the aggregator as an argument. The jump problem can then be solved with the discrete stepper provided by `JumpProcesses.jl`, `SSAStepper`. The code for simulating the homogeneous Poisson process with `Direct` is reproduced in Listing 1.

Listing 1: Simulation of the homogeneous Poisson process.

```

505 using JumpProcesses
506 rate(u, p, t) = p[1]
507 affect!(integrator) = (integrator.u[1] += 1;
508     nothing)
509
510 jump = ConstantRateJump(rate, affect!)
511 u, tspan, p = [0.], (0., 200.), (0.25,)
512 dprob = DiscreteProblem(u, tspan, p)
513 jprob = JumpProblem(dprob, Direct(), jump;

```



```

514     dep_graph=[[1]]
515     sol = solve(jprob, SSAS stepper())

```

517 The simulation of a Hawkes process — see Subsection 6.2 for a
518 definition — requires a `VariableRateJump` along with the rate
519 bounds and the interval for which the rates are valid. Also, since
520 the Hawkes process is history dependent, we close the `rate` and
521 `affect!` function with a vector containing the history of events.
522 The code for simulating the Hawkes process is reproduced in List-
523 ing 2. Note that it is possible to simplify the computation of the
524 rate — see Subsection 6.2 —, but we keep the code here as close
525 as possible to its usual definition for illustration purposes.

Listing 2: Simulation of the Hawkes process.

```

526 using JumpProcesses
527 h = Float64[]
528 rate(u, p, t) = p[1] +
529     p[2]*sum(exp(-p[3]*(t-_t)) for _t in h; init=0)
530 lrate(u, p, t) = p[1]
531 urate = rate
532 rateinterval(u, p, t) = 1/(2*urate(u,p,t))
533 affect!(integrator) = (push!(h, integrator.t);
534     integrator.u[1] += 1; nothing)
535 jump = VariableRateJump(rate, affect!; lrate,
536     urate, rateinterval)
537 u, tspan, p = [0.], (0., 200.), (0.25, 0.5, 2.0)
538 dprob = DiscreteProblem(u, tspan, p)
539 jprob = JumpProblem(dprob, Coevolve(), jump;
540     dep_graph=[[1]])
541 sol = solve(jprob, SSAS stepper())
542

```

544 To assess the correctness of `Coevolve`, we add it to the `Jump-`
545 `Processes.jl` test suite. Some tests check whether the aggrega-
546 tors are able to obtain empirical statistics close to the expected
547 in a number of simple biochemistry models such as linear reactions,
548 DNA repression, reversible binding and extinction. The test suite
549 was missing a unit test for self-exciting process. Thus, we have
550 added a test for the univariate Hawkes model that checks whether
551 algorithms that accept `VariableRateJump` are able to produce
552 an empirical distribution of trajectories whose first two moments of
553 the observed rate are close to the expected ones.


554 In addition to that, the correctness of the implemented algorithm
555 can be visually assessed using a Q-Q plot. As discussed in Sub-
556 section 4.1, every simple point process can be transformed to a
557 Poisson process with unit rate. This implies that the interval be-
558 tween points for any such transformed process should match the
559 exponential distribution. Therefore, the correctness of any aggrega-
560 tor can be assessed as following. First, transform the simulated
561 intervals with the appropriate compensator. Let t_{n_i} be the time in
562 which the n -th event of sub-process i took place and $t_{0_i} \equiv 0$, the
563 compensator for sub-process i is given by the following:

$$\Lambda_i^*(t_{n_i}) \equiv \Lambda_{n_i}^* \equiv \int_0^{t_{n_i}} \lambda_i^*(u) du \quad (6.1)$$


564 Then the transformed simulated interval is given by:

$$\Delta \Lambda_{n_i} \equiv \Lambda_{n_i}^* - \Lambda_{(n-1)_i}^* \quad (6.2)$$


565 Compute the empirical quantiles of the transformed intervals. That
566 is, the q -th quantile is the interval $\Delta \Lambda_q$ that divides the sorted
567 intervals in two sets, those below and above $\Delta \Lambda_q$ such that q -percent
568 of the elements are below it. Plot the empirical quantiles with the
569 corresponding quantiles of the exponential distribution. If the simu-
570 lator produces correct trajectories, this plot known as Q-Q plot




(a)



(b)



(c)



(d)

Fig. 1: Simulations of 10-nodes compound Hawkes process with parameters $\lambda = 0.5, \alpha = 0.1, \beta = 2.0$ for 200 units of time. (a) and (b) sampled trajectory and intensity rate for a single simulation for the three selected nodes in (c) for the first 20 units of time. (c) underlying 10-nodes network with three random nodes selected. (d) Q-Q plot of transformed inter-event time for 250 simulations colored by node.

571 should depict the points aligned around the 45-degree line. We
572 produce Q-Q plots for the homogeneous Poisson process as well as the
573 compound Hawkes process — see Subsection 6.2 for a definition
574 — to attest the correctness of `Coevolve`. Figure 1 (d) depicts the
575 Q-Q plot for a ten-node compound Hawkes process with paramet-
576 ers $\lambda = 0.5, \alpha = 0.1, \beta = 2.0$ simulated 250 times for 200
577 units of time. Figure 1 also depicts the trajectory, the conditional
578 intensity and the network structure of a single simulation for three
579 random nodes in panels (a), (b) and (c) respectively. We obtained
580 similar Q-Q plots for the other algorithms that benchmarked the
581 Multivariate Hawkes process below.

582 6.2 Benchmarks

583 We conduct a set of benchmarks to assess the performance of the
584 `JumpProcesses.jl` aggregators described in Section 5. All

| | Diffusion | Multi-state | Gene I | Gene II |
|---------------|---------------|---------------|----------------|---------------|
| Direct | 7.18 s | 0.16 s | 0.24 ms | 0.59 s |
| FRM | 15.04 s | 0.25 s | 0.29 ms | 0.78 s |
| SortingDirect | 1.08 s | <u>0.11 s</u> | 0.23 ms | 0.50 s |
| NRM | 0.75 s | 0.25 s | 0.39 ms | 0.89 s |
| DirectCR | <u>0.51 s</u> | 0.21 s | 0.47 ms | 1.00 s |
| RSSA | 1.42 s | 0.10 s | 0.43 ms | 0.65 s |
| RSSACR | 0.46 s | 0.16 s | 0.91 ms | 1.07 s |
| Coevolve | 0.90 s | 0.36 s | 0.59 ms | 1.33 s |

Table 1. : Median execution time. A 1-dimensional continuous time random walk approximation of a diffusion model (Diffusion), the multi-state model from Appendix A.6 [13] (Multi-state), a simple negative feedback gene expression model (Gene I) and the negative feedback gene expression from [7] (Gene II). Fastest time is **bold**, second fastest underlined. Benchmark source code and dependencies are available in SciMLBenchmarks.jl, see first paragraph of Section 6.2 for source references.

benchmarks are available in SciMLBenchmarks.jl³. All were run in BuildKite⁴ via the continuous integration facilities provided by the package maintainers. We have added two benchmark suites to assess the performance of the new aggregators under settings that could not be simulated with previous aggregators.

First, we assess the speed of the aggregators against jump processes whose rates are constant between jumps. There are four such benchmarks: a 1-dimensional continuous time random walk approximation of a diffusion model (Diffusion), the multi-state model from Appendix A.6 [13] (Multi-state), a simple negative feedback gene expression model (Gene I) and the negative feedback gene expression from [7] (Gene II). We simulate a single trajectory for each aggregator to visually check that they produce similar trajectories for a given model. The Diffusion, Multi-state, Gene I and Gene II benchmarks are then simulated 50, 100, 2000 and 200 times, respectively. Check the source code for further implementation details.

Benchmark results are listed in Table 1. The table shows that no single aggregator dominates suggesting they should be selected according to the task at hand. However, FRM, NRM, Coevolve never dominate any benchmark. In common, they all belong to the family of queuing methods suggesting that there is a penalty when using such methods for jump processes whose rates are constant between jumps. We also note that the performance of Coevolve lag that of NRM despite the fact that Coevolve should take the same number of steps as NRM when no VariableRateJump is used. The reason behind this discrepancy is likely due to implementation differences, but left for future investigation.

Second, we add a new benchmark which simulates the compound Hawkes process for an increasing number processes. Consider a graph with V nodes. The compound Hawkes process is characterized by V point processes such that the conditional intensity rate of node i connected to a set of nodes E_i in the graph is given by

$$\lambda_i^*(t) = \lambda + \sum_{j \in E_i} \sum_{t_{n_j} < t} \alpha \exp[-\beta(t - t_{n_j})]. \quad (6.3)$$

This process is known as self-exciting, because the occurrence of an event j at t_{n_j} will increase the conditional intensity of all the

processes connected to it by α . The excited intensity then decreases at a rate proportional to β .

$$\begin{aligned} \frac{d\lambda_i^*(t)}{dt} &= -\beta \sum_{j \in E_i} \sum_{t_{n_j} < t} \alpha \exp[-\beta(t - t_{n_j})] \\ &= -\beta(\lambda_i^*(t) - \lambda) \end{aligned} \quad (6.4)$$

The conditional intensity of this process has a recursive formulation which can significantly speed the simulation. The recursive formulation for the univariate case is derived in [10] which also provides additional discussion and results on the Hawkes process. We derive the compound case here. Let $t_{N_i} = \max\{t_{n_j} < t \mid j \in E_i\}$ and $\phi_i^*(t)$ below.

$$\begin{aligned} \phi_i^*(t) &= \sum_{j \in E_i} \sum_{t_{n_j} < t} \alpha \exp[-\beta(t - t_{N_i} + t_{N_i} - t_{n_j})] \\ &= \exp[-\beta(t - t_{N_i})] \sum_{j \in E_i} \sum_{t_{n_j} \leq t_{N_i}} \alpha \exp[-\beta(t_{N_i} - t_{n_j})] \\ &= \exp[-\beta(t - t_{N_i})] (\alpha + \phi_i^*(t_{N_i})) \end{aligned} \quad (6.5)$$

Then the conditional intensity can be re-written in terms of $\phi_i^*(t_{N_i})$.

$$\lambda_i^*(t) = \lambda + \phi_i^*(t) = \lambda + \exp[-\beta(t - t_{N_i})] (\alpha + \phi_i^*(t_{N_i})) \quad (6.6)$$

A random graph is sampled from the Erdős-Rényi model. This model assumes the probability of an edge between two nodes is independent of other edges, which we fix at 0.2. Note that this setup implies an increasing expected node degree.

We fix the Hawkes parameters at $\lambda = 0.5, \alpha = 0.1, \beta = 5.0$ ensuring the process does not explode and simulate models in the range from 1 to 95 nodes for 25 units of time. We simulate 50 trajectories with a limit of ten seconds to complete execution. For this benchmark, we save the state of the system exactly after each jump.

We assess the benchmark in eight different settings. First, we run the *inverse* method, *Coevolve* and *CHV simple* using the brute force formula of the intensity rate which loops through the whole history of past events — Equation 6.3. Second, we simulate the same three methods with the recursive formula — Equation 6.6. Next, we run the benchmark against *CHV full*. All *CHV* specifications are implemented with *PiecewiseDeterministicMarkovProcesses.jl*⁵ which is developed by Veltz, the author of the *CHV* algorithm discussed in Subsection 4.1. Finally, we run the benchmark using the Python library *Tick*⁶. This library implements a version of the thinning method for simulating the Hawkes process and implements a recursive algorithm for computing the intensity rate.

Table 2 shows that the *Inverse* method which relies on root finding is the most inefficient of all methods for any system size. For large system size this method is unable to complete all 50 simulation runs because it needs to find an ever larger number of roots of an ever larger system of differential equations.

The recursive implementation of the intensity rate brings a considerable boost to the simulations, placing *Coevolve* as one of the

³<https://github.com/SciML/SciMLBenchmarks.jl/tree/3bf650c1aae7b10e49cbd10e8f626d2a517f3e79/benchmarks/Jumps>

⁴<https://buildkite.com/julialang/scimlbenchmarks-dot-jl/builds/1326#01898802-ba51-4cd5-a31f-6c9b937b6146>

⁵<https://github.com/rveltz/PiecewiseDeterministicMarkovProcesses.jl>

⁶<https://github.com/X-DataInitiative/tick>

660 fastest algorithms. As shown in Algorithm 5, every sampled point
 661 in `Coevolve` requires a number of expected updates equal to the
 662 expected degree of the dependency graph. Therefore, it is able to
 663 complete non-exploding simulations efficiently.
 664 The Python library `Tick` remains competitive for smaller prob-
 665 lems, but gets considerably slower for bigger ones. Also, it is only
 666 specialized to the Hawkes process. Another drawback is that the
 667 library wraps the actual C++ implementation. In contrast, `Jump-`
 668 `Processes.jl` can simulate many other point processes with a
 669 relatively simple user-interface provided by the Julia language.
 670 There is substantial difference between the performance of recur-
 671 sive *CHV simple* and *CHV full*. The former does not make use
 672 of the derivative of the intensity function in Equation 6.4 which is
 673 more efficient to compute than the recursive rate in Equation 6.6.
 674 On the one hand, `Coevolve` clearly dominates *CHV simple*.
 675 On the other hand, *CHV full* is slower for smaller networks, but
 676 slightly faster than `Coevolve` for larger models. This change in
 677 relative performance occurs due to the rate of rejection in `Coe-`
 678 `volve` increasing in model size for this particular model. We com-
 679 pute the rejection rate as one minus the ratio between the number
 680 of jumps and the number of calls to the upper bound. A system
 681 with a single node sees a rejection rate of around 8 percent which
 682 rapidly increases to 80 percent when the system reaches 20 nodes
 683 and plateaus at around 95 percent with 95 nodes.
 684 Finally, we introduce a new benchmark which is intended to assess
 685 the performance of algorithms capable of simulating the stochastic
 686 model of hippocampal synaptic plasticity with geometrical read-
 687 out of enzyme dynamics proposed in [18]. For short, we denote it
 688 as the synapse model. We chose to benchmark this model as it is
 689 representative of a complex biochemical model. It couples a jump
 690 problem containing 98 jumps affecting 49 discrete variables with
 691 a stiff, ordinary differential equation problem containing 34 con-
 692 tinuous variables. Continuous variables affect jump rates while the
 693 discrete variables affect the continuous problem. There are 3 stages
 694 to the simulation: pre-synaptic evolution, glutamate release, and
 695 post-synaptic evolution. Among the algorithms considered, only
 696 the *inverse* method implemented in `JumpProcesses.jl`, `Coe-`
 697 `volve` and *CHV* are theoretically able to simulate the synapse
 698 model. However, in practice, only the last two complete at least
 699 one benchmark run. The original synapse problem was described
 700 as a piecewise deterministic Markov process, so we do not make
 701 the distinction between *CHV simple* and *full* in this benchmark.
 702 Benchmark results are displayed in Table 3. We observe that *CHV*
 703 is the fastest algorithm completing the synapse evolution in about
 704 half of the time it takes `Coevolve` with less than half of the allo-
 705 cations. Further investigation reveals that the thinning procedure in
 706 `Coevolve` reaches an average of 70 percent over all jumps which
 707 then leads to 2 to 3 times more function evaluations and Jaco-
 708 bians created compared to *CHV*. Our implementation adds stop-
 709 ping times via a call to `register_next_jump_time!` even for
 710 rejected jumps — we do not know a jump will be rejected until
 711 evaluated. This then leads the ODE solver to step to those times and
 712 make additional function evaluations and Jacobians. Lemaire *et*
 713 *al.* [11] performs a similar benchmark in which they compare the
 714 Hodgkin-Huxley model against different thinning conditions and
 715 against an ODE approximation. They find that a thinned algorithm
 716 with optimal boundary conditions can run significantly faster than
 717 the ODE approximation. Thus, there could be plenty of room to
 718 improve the performance of `Coevolve` in our setting.
 719 A disadvantage of *CHV* compared with `Coevolve` is that it sup-
 720 ports limited saving options by design. To save at pre-specified
 721 times would require using the fact that solutions are piecewise con-
 722 stant to determine solutions at times in-between jumps — and for

723 coupled ODE-jump problems would require root-finding to deter-
 724 mine when $s(u) = s_n$ for each desired saving time s_n in Equa-
 725 tion 4.8. The alternative proposed in [23] is to introduce an artificial
 726 jump to the model such as the homogeneous Poisson process with
 727 unit rate to sample the system at regular intervals. Alternatively,
 728 `Coevolve` allows saving at any arbitrary point. A common work-
 729 flow in simulating jump processes, particularly when interested in
 730 calculating statistics over time, is to pre-specify a precise set of
 731 times at which to save a simulation. In theory, this reduces mem-
 732 ory pressure, particularly for systems with large numbers of jumps,
 733 and can give increased computational performance relative to sav-
 734 ing the state at the occurrence of every jump. However, in the case
 735 of the synapse model, the number of candidate time rejections far
 736 surpasses the number of jumps. Therefore, reducing the number of
 737 saving points — *e.g.* only saving at start and end of the simulation
 738 — does not significantly reduce allocations or running time. Given
 739 these considerations, we decided to save after every jump and at
 740 regular pre-specified intervals that occur at the same frequency as
 741 the artificial saving jump used by *CHV*.

Another parameter that can affect the precision and speed of the
 742 synapse benchmark is the ODE solver. The author of `Piece-`
 743 `wiseDeterministicMarkovProcesses.jl` discuss some
 744 of these issues in Discourse⁷. Some ODE solvers can be faster and
 745 more precise. Due to time constraints, we have not investigated this
 746 matter. The synapse benchmark uses the `AutoTsit5(Rosen-`
 747 `brock23())` solver in both `Coevolve` and *CHV*. Further inves-
 748 tigation of this matter is left to future research.

750 7. Conclusion

751 This paper demonstrates that `JumpProcesses.jl` is a fast,
 752 general-purpose library for simulating evolutionary point pro-
 753 cesses. With the addition of `Coevolve`, any point process on the
 754 real line with a non-negative, left-continuous, history-adapted and
 755 locally bounded intensity rate can be simulated with this library.
 756 The objective of this paper was to bridge the gap between the treat-
 757 ment of point process simulation in statistics and biochemistry. We
 758 demonstrated that many of the algorithms developed in biochem-
 759 istry which served as the basis for the `JumpProcesses.jl` ag-
 760 gregators can be mapped to three general methods developed in
 761 statistics for simulating evolutionary point processes. We showed
 762 that the existing aggregators mainly differ in how they update and
 763 sample from the intensity rate and mark distribution. As we per-
 764 formed this exercise, we noticed the lack of an efficient aggregator
 765 for variable intensity rates in `JumpProcesses.jl`, a gap which
 766 `Coevolve` is meant to fill.

767 `Coevolve` borrows many enhancements from other aggregators
 768 in `JumpProcesses.jl`. However, there are still a number of
 769 ways forward. First, given the performance of the *CHV* algo-
 770 rithm in our benchmarks, we should consider adding it to `Jump-`
 771 `Processes.jl` as another aggregator so that it can benefit from
 772 tighter integration with the SciML organization and libraries. The
 773 saving behavior of *CHV* might pose a challenge when bringing
 774 this algorithm to the library. We could leverage the connection be-
 775 tween *inverse* and *thinning* methods illustrated in Subsection 4.2
 776 to attempt to develop a version of this algorithm that can evolve in
 777 synchrony with model time. Second, the new aggregator depends
 778 on the user providing bounds on the jump rates as well as the du-
 779 ration of their validity. In practice, it can be difficult to determine

⁷<https://discourse.julialang.org/t/help-me-beat-lsoda/88236>

| | V | Inverse | Brute Force | | Inverse | Coevolve | Recursive | | Tick |
|------|----|-------------------------|------------------------------|-------------------------|----------------------|--------------------------------|---------------|--------------------------------|--------------------------------|
| | | | Coevolve | CHV simple | | | CHV simple | CHV full | |
| Time | 1 | 113.7 μ s | 4.8 μs | 174.2 μ s | 112.1 μ s | <u>5.1 μs</u> | 175.6 μ s | 173.1 μ s | 31.4 μ s |
| | 10 | 17.5 ms | 211.8 μ s | 4.8 ms | 11.0 ms | 76.1 μs | 432.4 μ s | 579.0 μ s | <u>179.0 μs</u> |
| | 20 | 139.1 ms | 1.5 ms | 50.7 ms | 59.3 ms | 282.9 μs | 924.7 μ s | <u>884.4 μs</u> | 1.2 ms |
| | 30 | 415.3 ms <i>n=25</i> | 3.3 ms | 133.0 ms | 200.0 ms | 516.9 μs | 1.7 ms | <u>1.3 ms</u> | 3.7 ms |
| | 40 | 2.2 s <i>n=5</i> | 8.2 ms | 342.0 ms <i>n=30</i> | 1.6 s <i>n=7</i> | 1.0 ms | 2.5 ms | <u>1.6 ms</u> | 9.2 ms |
| | 50 | 5.1 s <i>n=2</i> | 16.9 ms | 722.0 ms <i>n=14</i> | 3.4 s <i>n=3</i> | 1.6 ms | 3.7 ms | <u>2.0 ms</u> | 21.2 ms |
| | 60 | 8.5 s <i>n=2</i> | 37.7 ms | 1.3 s <i>n=8</i> | 6.2 s <i>n=2</i> | 2.3 ms | 5.1 ms | <u>2.5 ms</u> | 45.0 ms |
| | 70 | 14.2 s <i>n=1</i> | 59.5 ms | 2.1 s <i>n=5</i> | 10.9 s <i>n=1</i> | <u>3.3 ms</u> | 6.8 ms | 3.0 ms | 87.5 ms |
| | 80 | 22.2 s <i>n=1</i> | 88.3 ms | 3.3 s <i>n=3</i> | 15.2 s <i>n=1</i> | <u>4.2 ms</u> | 9.0 ms | 3.3 ms | 142.2 ms |
| | 90 | 35.8 s <i>n=1</i> | 139.7 ms | 6.2 s <i>n=2</i> | 24.6 s <i>n=1</i> | <u>5.5 ms</u> | 11.9 ms | 3.8 ms | 241.9 ms |

Table 2. : Median execution time for the compound Hawkes process, V is the number of nodes and n is the total number of successful executions under ten seconds. Brute force refers to the implementation of the intensity rate looping through the whole history of past events. Recursive refers to a recursive implementation that only requires looking at the previous state of each node. Inverse and Coevolve are algorithms from JumpProcesses.jl, CHV is an algorithm from PiecewiseDeterministicMarkovProcesses.jl. See Subsection 4.1 for the distinction between CHV simple and CHV full. Tick is a Python library. All simulations were run 50 times except when stated otherwise under the running time. Fastest time is **bold**, second fastest underlined. Benchmark source code and dependencies are available in SciMLBenchmarks.jl, see first paragraph of Section 6.2 for source references.

| | Time | Allocation |
|----------|--------------|----------------|
| Inverse | - | - |
| Coevolve | <u>4.9 s</u> | <u>95.2 Mb</u> |
| CHV | 2.4 s | 43.8 Mb |

Table 3. : Median execution time and memory allocation. All simulations were run 50 times, a dash indicates that no runs were successful. Fastest time is **bold**, second fastest underlined. Benchmark source code and dependencies are available in SciMLBenchmarks.jl, see first paragraph of Section 6.2 for source references.

799 can be expressed as a branching process. There are simulation algo-
 800 rithms that already take advantage of this structure to leap through
 801 time [10]. It would be important to adapt these algorithms for gen-
 802 eral, compound branching processes to cater for a larger number of
 803 settings. Finally, JumpProcesses.jl also includes algorithms
 804 for jumps over two-dimensional spaces. It might be worth conduct-
 805 ing a similar comparative exercise to identify algorithms in statist-
 806 ics for 2- and N -dimensional processes that could also be added
 807 to JumpProcess.jl as it has the potential to become the go-to
 808 library for general point process simulation.

8. Acknowledgements

This project has been made possible in part by grant number 2021-237457 from the Chan Zuckerberg Initiative DAF, an advised fund of Silicon Valley Community Foundation. SAI was also partially supported by NSF-DMS 1902854.

9. References

- [1] Daryl J. Daley and David Vere-Jones. *An Introduction to the Theory of Point Processes: Volume I: Elementary Theory and Methods*. Probability and Its Applications, An Introduction to the Theory of Point Processes. Springer-Verlag, 2 edition. doi:10.1007/b97277.
- [2] Mehrdad Farajtabar, Yichen Wang, Manuel Gomez-Rodriguez, Shuang Li, Hongyuan Zha, and Le Song. COEVOLVE: A joint point process model for information diffusion and network evolution. 18(1). doi:10.5555/3122009.3122050.

780 these bounds a priori, particularly for models with many ODE vari-
 781 ables. Moreover, determining such bounds from an analytical solu-
 782 tion or the underlying ODEs does not guarantee their holding for
 783 the numerically computed solution (which is obtained via an ODE
 784 discretization), and so modifications may be needed in practice. A
 785 possible improvement would be for JumpProcesses.jl to deter-
 786 mine these bounds automatically taking into account the deriva-
 787 tive of the rates. Deriving efficient bounds require not only knowl-
 788 edge of the problem and a good amount of analytical work, but also
 789 knowledge about the numerical integrator. At best, the algorithm
 790 can perform significantly slower if a suboptimal bound or interval
 791 is used, at worst it can return incorrect results if a bound is incorrect
 792 — *i.e.* it can be violated inside the calculated interval of validity.
 793 Third, JumpProcesses.jl would benefit from further develop-
 794 ment in inexact methods. At the moment, support is limited to pro-
 795 cesses with constant rates between jumps and the only solver avail-
 796 able SimpleTauLeaping does not support marks. Inexact meth-
 797 ods should allow for the simulation of longer periods of time when
 798 only an event count per time interval is required. Hawkes processes

- 825 [3] Michael A. Gibson and Jehoshua Bruck. Efficient Exact 885
 826 Stochastic Simulation of Chemical Systems with Many 886
 827 Species and Many Channels. 104(9). doi:10.1021/jp993732q. 887
 828 [4] Daniel T. Gillespie. Approximate accelerated stochas- 888
 829 tic simulation of chemically reacting systems. 115(4). 889
 830 doi:10.1063/1.1378322. 890
 831 [5] Daniel T. Gillespie. Exact stochastic simulation of coupled 891
 832 chemical reactions. 81(25). doi:10.1021/j100540a008. 892
 833 [6] Daniel T Gillespie. A general method for numerically simu- 893
 834 lating the stochastic time evolution of coupled chemical reac- 894
 835 tions. 22(4). doi:10.1016/0021-9991(76)90041-3. 895
 836 [7] Abhishekh Gupta and Pedro Mendes. An Overview of
 837 Network-Based and -Free Approaches for Stochastic Simu-
 838 lation of Biochemical Systems. 6(1). doi:10.3390/computa-
 839 tion6010009.
 840 [8] Petter Holme. Fast and principled simulations of the
 841 SIR model on temporal networks. 16(2). doi:10.1371/jour-
 842 nal.pone.0246961.
 843 [9] Günter Last and Mathew Penrose. *Lectures on the Poisson*
 844 *Process*. Cambridge University Press, 1st edition edition.
 845 [10] Patrick J. Laub, Young Lee, and Thomas Taimre. *The Ele-*
 846 *ments of Hawkes Processes*. Springer International Pub-
 847 lishing. doi:10.1007/978-3-030-84639-8.
 848 [11] Vincent Lemaire, Michèle Thieullen, and Nicolas Thomas.
 849 Exact Simulation of the Jump Times of a Class of Piecewise
 850 Deterministic Markov Processes. 75(3). doi:10.1007/s10915-
 851 017-0607-4.
 852 [12] P. A. W. Lewis and G. S. Shedler. Simulation of Nonhomo-
 853 geneous Poisson Processes with Log Linear Rate Function.
 854 63(3). doi:10.2307/2335727. jstor:2335727.
 855 [13] Luca Marchetti, Corrado Priami, and Vo Hong Thanh. *Sim-*
 856 *ulation Algorithms for Computational Systems Biol-*
 857 *ogy*. Texts in Theoretical Computer Science. An EATCS Se-
 858 ries. Springer International Publishing. doi:10.1007/978-3-
 859 319-63113-4.
 860 [14] James M. McCollum, Gregory D. Peterson, Chris D. Cox,
 861 Michael L. Simpson, and Nagiza F. Samatova. The sort-
 862 ing direct method for stochastic simulation of biochemi-
 863 cal systems with varying reaction execution behavior. 30(1).
 864 doi:10.1016/j.compbiolchem.2005.10.007.
 865 [15] James Meiss. *Differential Dynamical Systems, Re-*
 866 *vised Edition*. Mathematical Modeling and Computa-
 867 tion. Society for Industrial and Applied Mathematics.
 868 doi:10.1137/1.9781611974645.
 869 [16] Y. Ogata. On Lewis' simulation method for point processes.
 870 27(1). doi:10.1109/TIT.1981.1056305.
 871 [17] Christopher Rackauckas and Qing Nie. DifferentialEqua-
 872 tions.jl A Performant and Feature-Rich Ecosystem for Solv-
 873 ing Differential Equations in Julia. 5(1). doi:10.5334/jors.151.
 874 [18] Yuri E. Rodrigues, Cezar M. Tigaret, H el ene Marie, Cian
 875 O'Donnell, and Romain Veltz. A stochastic model of hip-
 876 pocampal synaptic plasticity with geometrical readout of en-
 877 zyme dynamics. doi:10.1101/2021.03.30.437703.
 878 [19] Howard Salis and Yiannis Kaznessis. Accurate hybrid
 879 stochastic simulation of a system of coupled chemical or bio-
 880 chemical reactions. 122(5). doi:10.1063/1.1835951.
 881 [20] Alexander Slepoy, Aidan P. Thompson, and Steven J. Plimp-
 882 ton. A constant-time kinetic Monte Carlo algorithm for sim-
 883 ulation of large biochemical reaction networks. 128(20).
 884 doi:10.1063/1.2919546.
 885 [21] Vo Hong Thanh, Corrado Priami, and Roberto Zunino.
 886 Efficient rejection-based simulation of biochemical
 887 reactions with stochastic noise and delays. 141(13).
 888 doi:10.1063/1.4896985.
 889 [22] Vo Hong Thanh, Roberto Zunino, and Corrado Pri-
 890 ami. Efficient Constant-Time Complexity Algorithm for
 891 Stochastic Simulation of Large Reaction Networks. 14(3).
 892 doi:10.1109/TCBB.2016.2530066.
 893 [23] Romain Veltz. A new twist for the simulation
 894 of hybrid systems using the true jump method.
 895 doi:10.48550/arXiv.1504.06873. arxiv:1504.06873.